

NYC Yellow Taxi Trip Records


Insights and Trends



Aiden Pham Dang Khoa
8 June 2026

OUTLINE

1. Introduction
2. Data quality
3. Data cleaning
4. Feature engineering
5. Data analysis
6. Conclusions



OUTLINE

1. Introduction

1.1 Dataset

1.2 Research questions

1.3 Business significance

1.4 Methods


2. Data quality

3. Data cleaning

4. Feature engineering

5. Data analysis

6. Conclusions



1. INTRODUCTION

1.1 Dataset


- ~**305,000** NYC Yellow Taxi trips from **January 2018 to January 2023**
- 2018–2022: ~**60,000** trips per year; 2023: January only (~**5,000** trips)
- Data appears sample-capped at ~5,000 trips/month - ratio metrics are reliable, absolute volumes are not.
- **19 columns:**
 - Pickup & dropoff time
 - Pickup & dropoff zone ID
 - Trip distance
 - Other dimensions (vendor, rate code, payment type, etc.)
 - Financial columns (fare, tax, surcharge, total, etc.)

RangeIndex: 304978 entries, 0 to 304977

Data columns (total 19 columns):

#	Column	Non-Null Count	Dtype
0	VendorID	304978 non-null	category
1	tpep_pickup_datetime	304978 non-null	datetime64[ns]
2	tpep_dropoff_datetime	304978 non-null	datetime64[ns]
3	passenger_count	295465 non-null	category
4	trip_distance	304978 non-null	float64
5	RatecodeID	295465 non-null	category
6	store_and_fwd_flag	295465 non-null	category
7	PULocationID	304978 non-null	category
8	DOLocationID	304978 non-null	category
9	payment_type	304978 non-null	category
10	fare_amount	304978 non-null	float64
11	extra	304978 non-null	float64
12	mta_tax	304978 non-null	float64
13	tip_amount	304978 non-null	float64
14	tolls_amount	304978 non-null	float64
15	improvement_surcharge	304978 non-null	float64
16	total_amount	304978 non-null	float64
17	congestion_surcharge	232346 non-null	float64
18	airport_fee	106217 non-null	float64

dtypes: category(7), datetime64[ns](2), float64(10)



1. INTRODUCTION

1.2 Research Questions

Analyzed in section	Topics
5.1 Overview	Citywide demand, revenue, payment, and efficiency patterns
5.2 Time	Demand and efficiency by month, weekday, and hour
5.3 Space	Demand and efficiency hotspots by zone and borough
5.4 Vendor	Vendor performance comparison
5.5 Rate code	Rate-code segment performance comparison
5.6 Tip	Tip behavior by time and location
5.7 Fare	Relationship between fare, distance, and duration



1. INTRODUCTION


1.3 Business significance

- Provide **clear, actionable KPIs**.
- Support **smarter fleet allocation**.
- Optimize **revenue management**.
- Enable more **targeted service strategy**.
- Improve **data quality**.



1. INTRODUCTION

1.4 Method



2. DATA QUALITY



In the scope of this assignment, I will focus on Completeness, Uniqueness, Validity and Accuracy.

The other 2 dimensions are not relevant in this assignment:

- Timeliness (how up-to-date data are): this is a static dataset
- Consistency (data are consistent among tables, systems): there is only 1 dataset

2. DATA QUALITY

2.1 Completeness

- ~3% of rows had NULL values in passenger_count, RatecodeID, and store_and_fwd_flag — all three missing together, suggesting a systematic data collection issue
- congestion_surcharge and airport_fee have NULLs that are not substitutes for zero — retained as missing

2.2 Uniqueness

- 3 pairs of duplicate trips identified based on identical pickup/dropoff time and distance
- Each pair showed a specific issue: NULL metadata (2 pairs) or negative financial values (1 pair)

2.3 Validity

- VendorID contains undocumented values (4 and 5) - treated as “Others”
- Negative values present in all numeric columns (less than 1%); zero-distance trips peaked in 2020
- Store-and-forward trips exhibit a substantially higher rate of zero-distance records (4.41%) compared with non-store-and-forward trips (1.20%)

2.4 Accuracy

- Implausible trips found: >100 miles, >3 hours duration, >100 mph speed
- total_amount mismatches linked to congestion_surcharge recording errors (confirmed by chi-square test)

3. DATA CLEANING

3.1 Completeness: NULL imputation

- passenger_count, RatecodeID, store_and_fwd_flag: mode-imputed with a missingness flag column added
- congestion_surcharge and airport_fee: retained as NULL (not forced to 0)

3.2 Uniqueness: duplicate removal

- Deleted the NULL-metadata row from each of the 2 metadata duplicate pairs
- Deleted the negative-value row from the 1 financial duplicate pair

3.3 Validity: remove negative values

- Dropped rows with negative trip_distance and financial columns
- Zero values retained — removing them silently would undercount minimum-fare and cancellation revenue

3.4 Accuracy: plausibility filters

- Kept only trips with: distance ≤ 100 miles, duration ≤ 180 minutes, speed ≤ 100 mph
- total_amount mismatches not corrected - no reliable business rule to determine which field is wrong

Result

- 0.82% of rows removed; cleaned dataset retains over 302k trips with more reasonable distributions

4. FEATURE ENGINEERING

4.1 New datetime and derived columns

- Pickup year, month, day-of-week, hour of day - for temporal analysis
- trip_duration_min, avg_speed_mph - for accuracy checks and efficiency analysis
- Human-readable labels for VendorID, RatecodeID, payment_type

4.2 Geospatial enrichment

- Joined TLC zone lookup table: added pickup/dropoff borough, zone name, and service zone
- Loaded TLC zone shapefiles for geographic heatmap visualizations


4.3 Sample Validation

Representativeness:

- Tested payment mix, borough share, and hourly distribution stability across months
- Sample composition is internally consistent; ratio metrics are suitable for trend analysis

COVID Impact:

- Many metrics experienced disruption in 2020
- All Correlation with COVID-19 noted; causality cannot be established from this sample alone



OUTLINE

1. Introduction

2. Data quality

3. Data cleaning

4. Feature engineering

5. Data analysis

5.1 L0 High-level overview

5.2 L1 Time

5.3 L1 Space

5.4 L1 Vendor

5.5 L1 Rate code



5.6 L2 Tip

5.7 L2 Fare

6. Conclusions




5. DATA ANALYSIS



On the left is my general data analysis framework. In this assignment, I used the modified framework (on the right).

5. DATA ANALYSIS

5.1 L0 High-level Overview



We will analyze the above 4 groups of metrics in this section.

5. DATA ANALYSIS



5.1 L0 High-level Overview

a) Demand



There is disruption in demand around March 2020 (will be discussed in 5.2).

The hourly profile indicates clear late-afternoon and early-evening intensity, while day-of-week results show stronger demand on weekdays than weekends.


Demand is the highest in Manhattan (about 90% of total number of trips).



Pickup Demand by Borough



Dropoff Demand by Borough




5. DATA ANALYSIS

5.1 L0 High-level Overview

b) Revenue

- Average revenue per trip trends upward and peaks in 2023.
- Revenue composition remains fare-led, with tip amount as the second-largest contributor.
- Congestion surcharge seems to be introduced in 2019.




5. DATA ANALYSIS

5.1 L0 High-level Overview

c) Payment & Tip

- Card payments dominate across all years and increase over time, while cash share generally declines.
- The share of tipped trips remained stable at around 95%.
- Average tip amount also trends upward (roughly \$2.9 to \$4.3), with a dip in 2020 and strong growth afterward.
- Overall, payment behavior is becoming more cash-less and tip amount is improving over time.



5. DATA ANALYSIS

5.1 L0 High-level Overview

d) Efficiency



- Average trip distance stays fairly stable across years, while average trip duration slightly dipped in 2020.
- The heatmap shows clear time-dependency, with faster speeds after midnight and slower speeds during office hours on weekdays.
- Revenue per mile and per minute jumped significantly in 2023.



5. DATA ANALYSIS

5.2 L1 Time


- The heatmap on the left shows that there is lower demand on the holidays (4 July and Christmas week). March seems highly concentrated in the early days of the month.
- The heatmap on the right explains that 2020 is the outlier, which correlates with Covid timing when there was no demand in late March 2020.



5. DATA ANALYSIS

5.2 L1 Time

- Trip distance is highest in late-night to early-morning windows, especially around 3-6 AM.
- In contrast, Revenue per mile is weaker in overnight periods.
- Revenue per minute has opposite pattern to Revenue per mile.
- Together, the three heatmaps suggest distance influence revenue efficiency. Overnight periods have longer-distance trips and better Revenue per minute efficiency, but lower Revenue per mile.




5. DATA ANALYSIS


5.3 L1 Space

- The trips are **highly concentrated** in a few zones: out of ~250 zones, both the top 20 pickup and dropoff zones account for more than half of trip volume.
- The top 20 pickup zones include 18 in Manhattan and 2 airports in Queens.
- The top 20 dropoff zones are all in Manhattan.
- Nearly 85% of trips start and end within Manhattan, which is expected given that Manhattan is the most densely populated borough and a major commercial hub.

Share of Trips in Top 20 Pickup and Dropoff Zones




Sankey Chart: Pickup Borough to Dropoff Borough (colored by Pickup Borough)



5. DATA ANALYSIS

5.3 L1 Space





- **Trip volume** is heavily concentrated in Manhattan, with airport-related zones in Queens as secondary hotspots.
- **Average trip distance** is generally shorter in high-volume core zones (especially Manhattan) and longer in lower-volume outer-borough zones.
- **Revenue per mile** tends to be strongest in central, dense activity zones and weaker across many peripheral zones.
- **Revenue per minute** adds a complementary view: the general pattern is opposite to Revenue per mile.
- Together, the four maps show a clear core-periphery pattern: a high-demand urban core with stronger distance-based yield versus a lower-demand outer area with more mixed efficiency outcomes.



5. DATA ANALYSIS

5.4 L1 Vendor





- Trip volume is highly concentrated in Curb Mobility and Creative Mobile Technologies; other vendors have minimal share.
- The two major vendors maintain stable trip distance, while smaller vendors tend to run longer trips.
- Revenue per mile and per minute is stronger for the top vendors and jumps in 2023.
- Myle remains lower for both efficiency metrics, suggesting a different fare mix.
- Results for small vendors are volatile and should be interpreted cautiously due to low trip counts.



5. DATA ANALYSIS

5.5 L1 Rate Code


- Standard rides drive the business by volume, while non-standard codes remain niche.
- Airport and codes carry longer trip profiles.
- Negotiated rate code perform the best for both revenue efficiency metrics.



5. DATA ANALYSIS

5.6 L2 Tip



- Tips are right-skewed: most trips have low-to-moderate tips, while high tips are relatively rare.
- The concentration near small tip values suggests many short trips or low-fare rides. A thin long tail indicates occasional high-tip trips, likely linked to longer distances.
- A large majority of trips have 20% tip. This may be the default choice for customers.
- Many trips have tip amount at \$1 and \$2.
- Positive correlation between tip amount vs trip distance and trip duration (Spearman coefficient ~ 0.68).



5. DATA ANALYSIS

5.6 L2 Tip




- Zones with the highest tips are mostly in peripheral areas and the airports.
- Tip amounts tend to be higher after midnights.



5. DATA ANALYSIS

5.7 L2 Fare



- The distribution of fare amounts is right-skewed, with a long tail of higher fares.
- Fare rises with both trip distance and trip duration, confirmed by strong positive relationships in both scatter plots and Spearman test (coefficient ~ 0.9).
- The fare-distance pattern is tighter and more linear than fare-duration, suggesting distance is the cleaner direct driver of pricing.
- The fare-duration plot is more dispersed at the top right, indicating that other factors introduce additional fare variability beyond travel time alone.



5. DATA ANALYSIS

5.7 L2 Fare

- Manhattan zones generally achieve higher fare per mile but lower fare per minute, consistent with shorter, slower trips in congested urban areas.
- Temporal patterns suggest congestion increases fare per mile during peak periods, while off-peak periods tend to produce more time-efficient trips and higher fare per minute.
- The contrasting behavior of these 2 metrics supports a zone- and time-specific operating model, with both KPIs used jointly for operational monitoring and shift planning.



6. CONCLUSIONS

NYC TAXI DATASET – KEY FINDINGS



1 TIME & LOCATION MATTER MOST



- Performance is highly dependent on time and location.
- Citywide averages mask significant zone-and time-level differences.
- Decisions should be informed at the zone and time level.

2 WEEKDAYS LATE AFTERNOONS, AND MANHATTAN DOMINATE DEMAND



- Manhattan accounts for the majority of trips.
- Weekdays see higher demand.
- Late afternoons are the busiest time of day.

3 REVENUE IMPROVED OVER TIME, ESPECIALLY IN 2023



- Total revenue shows a clear upward trend over the years.
- The strongest improvement was observed in 2023.

4 CARD ACCOUNTS FOR OVER 75% OF PAYMENT. MOST TRIPS HAVE 20% TIPPING RATE.



- Card payments dominate transactions.
- Common tipping rate is 20%.

5 EFFICIENCY PER MILE AND PER MINUTE CAN MOVE IN DIFFERENT DIRECTION



- Revenue per mile and revenue per minute do not always move together.
- Management decisions should balance both efficiency metrics.

6 DEMAND IS CONCENTRATED IN TWO VENDORS AND STANDARD-RATE TRIPS



- A small number of vendors account for the majority of trips.
- Standard-rate trips represent the largest share of demand.

7 FARE AND TIPS INCREASE WITH TRIP DISTANCE AND DURATION



- Fare and Tip are higher for trips with longer distance and longer duration.
- Both are different per time and location.



The NYC taxi market is driven by time, location, and trip characteristics, with demand concentrated in Manhattan, card payments dominating transactions, and revenue improving over time. Operational decisions should focus on zone-level performance and balance both revenue-per-mile and revenue-per-minute efficiency metrics.



6. CONCLUSIONS

BUSINESS RECOMMENDATIONS



1



Adopt a Zone- and Time-Specific Operating Model

Monitor KPIs at the zone and time-segment level to capture meaningful differences.
Use insights to drive targeted driver allocation, shift planning, and performance management.

2



Use Revenue per Mile and Revenue per Minute as Complementary KPIs

These metrics highlight different operational strengths—short urban trips excel in revenue per mile, while longer peripheral trips perform better in revenue per minute.
Consider both for balanced, data-driven decisions.

3



Prioritize Peak Demand Periods for Capacity Planning

Trip volume follows clear daily and weekly patterns.
Align driver availability with peak demand to improve service and reduce underutilization.

4



Investigate High-Efficiency Route and Rate-Code Segments

Some zones and rate codes consistently outperform others.
Evaluate their scalability and the volume behind the performance before prioritizing for optimization.

5



Establish Ongoing Data-Quality Monitoring

Detect and address missing or unrealistic records automatically.
Strengthen data quality to ensure reliable metrics and decisions.





Questions?